

ommerce-marketing-and-shopping-eda

August 29, 2024

```
[1]: import pandas as pd
import numpy as np
from scipy import stats
import matplotlib.pyplot as plt
import seaborn as sns
from sklearn.preprocessing import StandardScaler
import scipy.stats as stats
from scipy.stats import kruskal
from datetime import datetime
```

```
[2]: shop_df = pd.read_csv('/content/shopping.csv')
```

```
[3]: shop_df.head()
```

```
[3]:
```

	Administrative	Administrative_Duration	Informational	\
0	0	0.0	0	
1	0	0.0	0	
2	0	0.0	0	
3	0	0.0	0	
4	0	0.0	0	

	Informational_Duration	ProductRelated	ProductRelated_Duration	\
0	0.0	1	0.000000	
1	0.0	2	64.000000	
2	0.0	1	0.000000	
3	0.0	2	2.666667	
4	0.0	10	627.500000	

	BounceRates	ExitRates	PageValues	SpecialDay	Month	OperatingSystems	\
0	0.20	0.20	0.0	0.0	Feb	1	
1	0.00	0.10	0.0	0.0	Feb	2	
2	0.20	0.20	0.0	0.0	Feb	4	
3	0.05	0.14	0.0	0.0	Feb	3	
4	0.02	0.05	0.0	0.0	Feb	3	

	Browser	Region	TrafficType	VisitorType	Weekend	Revenue
0	1	1	1	Returning_Visitor	False	False

1	2	1	2	Returning_Visitor	False	False
2	1	9	3	Returning_Visitor	False	False
3	2	2	4	Returning_Visitor	False	False
4	3	1	4	Returning_Visitor	True	False

[3]:

1 Data Preprocessing

2 1. Shopping dataframe

[4]: `shop_df.shape`

[4]: (12330, 18)

In the shopping dataset we have 12330 rows and 18 columns

[5]: `shop_df.info()`

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 12330 entries, 0 to 12329
Data columns (total 18 columns):
#   Column                                Non-Null Count  Dtype
---  -
0   Administrative                        12330 non-null  int64
1   Administrative_Duration              12330 non-null  float64
2   Informational                        12330 non-null  int64
3   Informational_Duration               12330 non-null  float64
4   ProductRelated                      12330 non-null  int64
5   ProductRelated_Duration              12330 non-null  float64
6   BounceRates                         12330 non-null  float64
7   ExitRates                          12330 non-null  float64
8   PageValues                         12330 non-null  float64
9   SpecialDay                         12330 non-null  float64
10  Month                              12330 non-null  object
11  OperatingSystems                   12330 non-null  int64
12  Browser                           12330 non-null  int64
13  Region                            12330 non-null  int64
14  TrafficType                       12330 non-null  int64
15  VisitorType                       12330 non-null  object
16  Weekend                           12330 non-null  bool
17  Revenue                           12330 non-null  bool
dtypes: bool(2), float64(7), int64(7), object(2)
memory usage: 1.5+ MB
```

[6]: `shop_df.isnull().sum()`

```
[6]: Administrative      0
      Administrative_Duration  0
      Informational      0
      Informational_Duration  0
      ProductRelated      0
      ProductRelated_Duration  0
      BounceRates      0
      ExitRates      0
      PageValues      0
      SpecialDay      0
      Month      0
      OperatingSystems      0
      Browser      0
      Region      0
      TrafficType      0
      VisitorType      0
      Weekend      0
      Revenue      0
      dtype: int64
```

We don't have any null values in the shopping dataset

```
[7]: shop_df[shop_df.duplicated()]
```

```
[7]:      Administrative  Administrative_Duration  Informational  \
158                0                0.0                0
159                0                0.0                0
178                0                0.0                0
418                0                0.0                0
456                0                0.0                0
...                ...                ...                ...
11934               0                0.0                0
11938               0                0.0                0
12159               0                0.0                0
12180               0                0.0                0
12185               0                0.0                0

      Informational_Duration  ProductRelated  ProductRelated_Duration  \
158                0.0                1                0.0
159                0.0                1                0.0
178                0.0                1                0.0
418                0.0                1                0.0
456                0.0                1                0.0
...                ...                ...                ...
11934               0.0                1                0.0
11938               0.0                1                0.0
12159               0.0                1                0.0
```

12180		0.0		1		0.0
12185		0.0		1		0.0

	BounceRates	ExitRates	PageValues	SpecialDay	Month	OperatingSystems	\
158	0.2	0.2	0.0	0.0	Feb		1
159	0.2	0.2	0.0	0.0	Feb		3
178	0.2	0.2	0.0	0.0	Feb		3
418	0.2	0.2	0.0	0.0	Mar		1
456	0.2	0.2	0.0	0.0	Mar		2
...
11934	0.2	0.2	0.0	0.0	Dec		1
11938	0.2	0.2	0.0	0.0	Dec		1
12159	0.2	0.2	0.0	0.0	Dec		1
12180	0.2	0.2	0.0	0.0	Dec		1
12185	0.2	0.2	0.0	0.0	Dec		8

	Browser	Region	TrafficType	VisitorType	Weekend	Revenue
158	1	1	3	Returning_Visitor	False	False
159	2	3	3	Returning_Visitor	False	False
178	2	3	3	Returning_Visitor	False	False
418	1	1	1	Returning_Visitor	True	False
456	2	4	1	Returning_Visitor	False	False
...
11934	1	1	2	New_Visitor	False	False
11938	1	4	1	Returning_Visitor	True	False
12159	1	1	3	Returning_Visitor	False	False
12180	13	9	20	Returning_Visitor	False	False
12185	13	9	20	Other	False	False

[125 rows x 18 columns]

```
[8]: shop_df.drop_duplicates(inplace=True)
```

```
[9]: shop_df.duplicated().sum()
```

```
[9]: 0
```

These are duplicated values in the shopping dataset. Removed it for clear analysis.

2.0.1 - outliers

```
[10]: numerical_col = shop_df.select_dtypes(include=np.number).columns
```

```
[11]: numerical_col
```

```
[11]: Index(['Administrative', 'Administrative_Duration', 'Informational',
          'Informational_Duration', 'ProductRelated', 'ProductRelated_Duration',
```

```

'BounceRates', 'ExitRates', 'PageValues', 'SpecialDay',
'OperatingSystems', 'Browser', 'Region', 'TrafficType'],
dtype='object')



```



Visualizing Outliers



```



[12]: for col in numerical_col:
plt.figure(figsize=(10,5))
sns.boxplot(shop_df[col])
plt.show()



```























Univeriant Analysis



```
[13]: for col in numerical_col:
      plt.figure(figsize=(10,5))
      sns.histplot(shop_df[col])
      plt.show()
```




















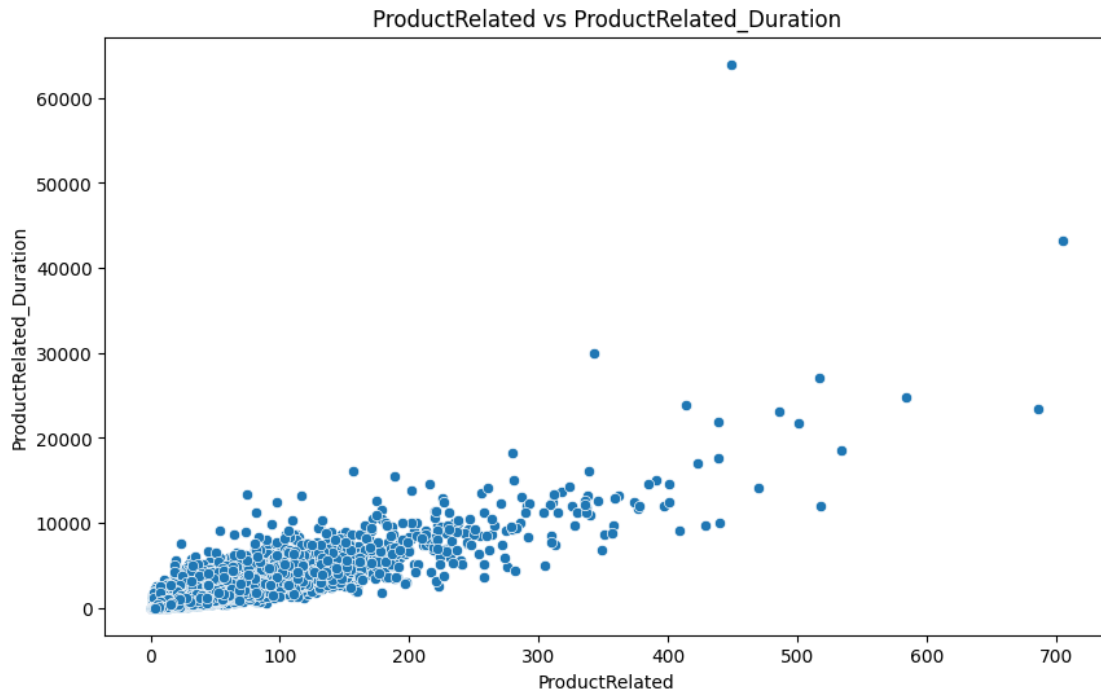
- Correlation Analysis

```
[14]: correlation_matrix = shop_df[numerical_col].corr()
plt.figure(figsize=(12, 10))
sns.heatmap(correlation_matrix, annot=True, cmap='coolwarm', vmin=-1, vmax=1)
plt.title('Correlation Matrix')
plt.show()
```



- Visualizations

```
[15]: plt.figure(figsize=(10, 6))
sns.scatterplot(x='ProductRelated', y='ProductRelated_Duration', data=shop_df)
plt.title('ProductRelated vs ProductRelated_Duration')
plt.xlabel('ProductRelated')
plt.ylabel('ProductRelated_Duration')
plt.show()
```




Users spending more time on product-related pages could be more likely to convert into customers. This suggests a need to optimize these pages to encourage further engagement and eventually drive conversions.

Class Distribution

```
[16]: revenue_distribution = shop_df['Revenue'].value_counts(normalize = True)
revenue_distribution
```

```
[16]: Revenue
False    0.843671
True     0.156329
Name: proportion, dtype: float64
```

```
[17]: plt.figure(figsize=(6, 4))
sns.countplot(x='Revenue', data=shop_df)
plt.title('Class Distribution of Revenue')
plt.show()
```



Summarize page views, durations, and bounce/exit rates for each page category.

```
[18]: shop_df.columns
```

```
[18]: Index(['Administrative', 'Administrative_Duration', 'Informational',
        'Informational_Duration', 'ProductRelated', 'ProductRelated_Duration',
        'BounceRates', 'ExitRates', 'PageValues', 'SpecialDay', 'Month',
        'OperatingSystems', 'Browser', 'Region', 'TrafficType', 'VisitorType',
        'Weekend', 'Revenue'],
        dtype='object')
```

```
[19]: page_summary = shop_df.groupby('ProductRelated').agg({
        'Administrative': 'sum',
        'Administrative_Duration': 'sum',
        'Informational': 'sum',
        'Informational_Duration': 'sum',
        'BounceRates': 'mean',
        'ExitRates': 'mean'
    })
print(page_summary)
```

```
Administrative  Administrative_Duration  Informational \
```

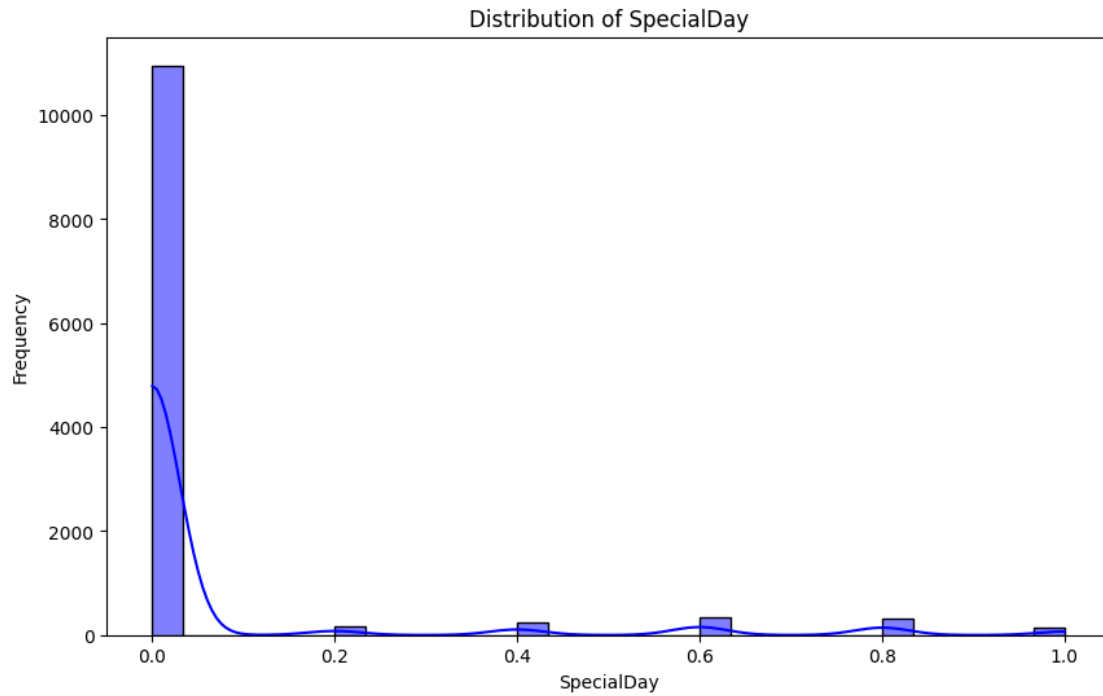
ProductRelated			
0	35	847.400000	12
1	67	2867.450000	31
2	232	8854.512500	30
3	379	13080.666667	23
4	332	12124.116349	64
...
518	8	161.668571	0
534	9	444.284722	0
584	27	853.735949	2
686	20	199.456273	7
705	17	2629.253968	24

	Informational_Duration	BounceRates	ExitRates
ProductRelated			
0	89.000000	0.098158	0.127446
1	1751.650000	0.166177	0.179688
2	1064.750000	0.046546	0.103480
3	643.400000	0.033990	0.077767
4	3723.666667	0.027736	0.064376
...
518	0.000000	0.000038	0.003837
534	0.000000	0.010857	0.023309
584	126.500000	0.002099	0.009347
686	299.033333	0.009853	0.022771
705	2050.433333	0.004851	0.015431

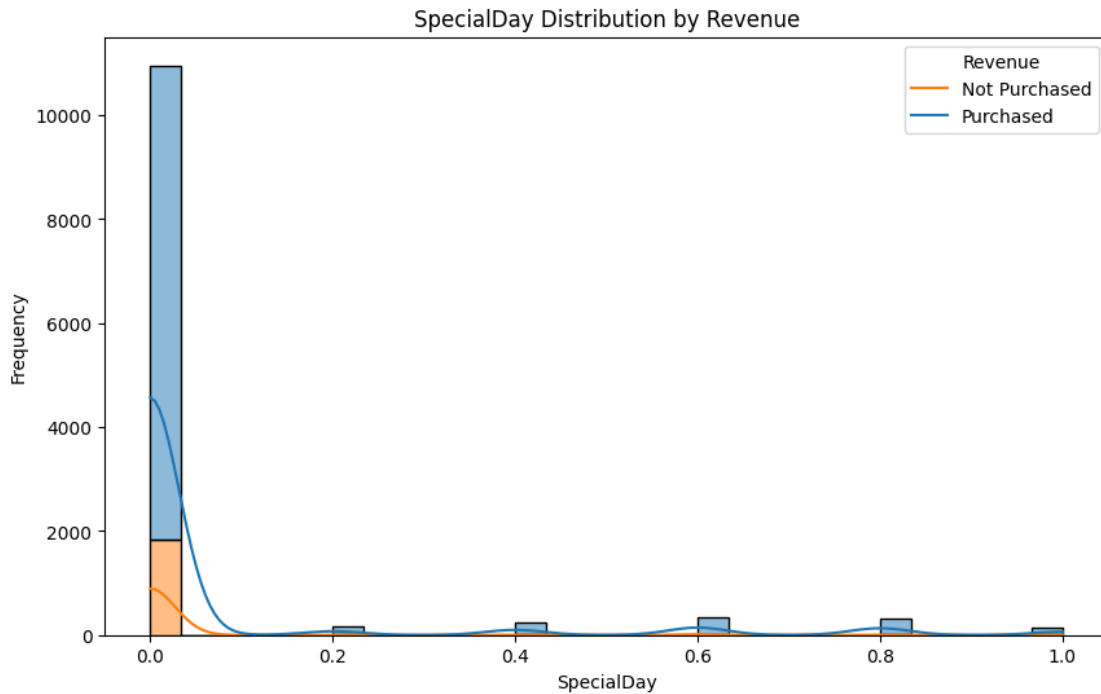
[311 rows x 6 columns]

Analyze SpecialDay Distribution and Its Correlation with Revenue

```
[20]: plt.figure(figsize=(10, 6))
sns.histplot(shop_df['SpecialDay'], bins=30, kde=True, color='blue')
plt.title('Distribution of SpecialDay')
plt.xlabel('SpecialDay')
plt.ylabel('Frequency')
plt.show()
```



```
[21]: plt.figure(figsize=(10, 6))
sns.histplot(data=shop_df, x='SpecialDay', hue='Revenue', multiple='stack',
             bins=30, kde=True)
plt.title('SpecialDay Distribution by Revenue')
plt.xlabel('SpecialDay')
plt.ylabel('Frequency')
plt.legend(title='Revenue', labels=['Not Purchased', 'Purchased'])
plt.show()
```



Generate a binary feature indicating whether the user visited all three page categories.

```
[22]: shop_df.columns
```

```
[22]: Index(['Administrative', 'Administrative_Duration', 'Informational',
            'Informational_Duration', 'ProductRelated', 'ProductRelated_Duration',
            'BounceRates', 'ExitRates', 'PageValues', 'SpecialDay', 'Month',
            'OperatingSystems', 'Browser', 'Region', 'TrafficType', 'VisitorType',
            'Weekend', 'Revenue'],
          dtype='object')
```

```
[23]: shop_df['All_Pages_Visited'] = ((shop_df['Administrative'] > 0) &
                                       (shop_df['Informational'] > 0) &
                                       (shop_df['ProductRelated'] > 0)).astype(int)
```

```
[24]: shop_df['All_Pages_Visited'].value_counts()
```


```
[24]: All_Pages_Visited
0      10038
1       2167
Name: count, dtype: int64
```

users are not exploring all content types, possibly missing critical information.

Recommendation: Improve Cross-Linking: Ensure that there are clear and intuitive links between different types of content. For instance, product pages could link to related informational content or FAQs.

Explore PageValues distribution and its relationship with TrafficType, VisitorType, and Region

```
[25]: sns.histplot(shop_df['PageValues'], kde=True)
plt.title('Distribution of PageValues')
plt.xlabel('PageValues')
plt.ylabel('Frequency')
plt.show()
```



- relationship between traffic and pagevalue

```
[26]: corr = shop_df[['PageValues', 'TrafficType']].corr()
corr
```

```
[26]:
```

	PageValues	TrafficType
PageValues	1.000000	0.012286


TrafficType 0.012286 1.000000

- relationship between visitorType and pagevalue

```
[27]: visitor_stats = shop_df.groupby('VisitorType')['PageValues'].agg(['mean', 'median', 'std', 'count'])
print(visitor_stats)
```

	mean	median	std	count
VisitorType				
New_Visitor	10.778550	0.0	29.197231	1693
Other	19.090173	0.0	54.328142	81
Returning_Visitor	5.063768	0.0	15.501615	10431

```
[28]: residuals = shop_df['PageValues']
stats.probplot(residuals, dist="norm", plot=plt)
plt.title('Q-Q Plot of Residuals / PageValues')
plt.show()
```




```
[29]: kruskal_test = kruskal(*[group['PageValues'].values for name, group in shop_df.
    ↳groupby('VisitorType')])
print("Kruskal-Wallis Test:", kruskal_test)
```

Kruskal-Wallis Test: KruskalResult(statistic=6.207841451955822, pvalue=0.04487292261362119)

- Relationship between pagevalues and region

```
[30]: pivot_table = shop_df.pivot_table(values='PageValues', index='Region',
    ↳columns='TrafficType', aggfunc='mean')
sns.heatmap(pivot_table, annot=True, cmap='coolwarm')
plt.title('Heatmap of PageValues by Region and TrafficType')
plt.show()
```



- Investigate user session lengths and their impact on conversion rates.

```
[31]: shop_df['sessionlength'] = shop_df['Administrative'] + shop_df['Informational']
    ↳+ shop_df['ProductRelated']
```

```
[32]: shop_df['sessionlength'].max()
```

[32]: 746

```
[33]: shop_df['Session_category'] = pd.cut(shop_df['sessionlength'],  
      ↪bins=[0,250,500,750,np.inf], labels = ['Short', 'Medium', 'Long', 'Very'  
      ↪Long'])
```

```
[34]: shop_df['Session_category'].value_counts()
```

```
[34]: Session_category  
Short      12106  
Medium      86  
Long        7  
Very Long   0  
Name: count, dtype: int64
```


Insight : Longer session lengths generally lead to higher conversion rates, with “Long” sessions having a conversion rate of 42.86%. However, “Very Long” sessions didn’t result in conversions.

Recommendation: Focus on optimizing the user experience to maintain engagement during “Long” sessions without overstaying, as this might lead to frustration or indecision. Tools like personalized recommendations or limited-time offers could help convert these sessions.

```
[35]: conversion_rate = shop_df.groupby('Session_category', observed =  
      ↪False)['Revenue'].mean() * 100  
print(conversion_rate)
```

```
Session_category  
Short      15.488188  
Medium     34.883721  
Long       42.857143  
Very Long   NaN  
Name: Revenue, dtype: float64
```

```
[36]: sns.scatterplot(x='sessionlength', y='Revenue', data=shop_df)  
plt.title('Session Length vs. Conversion')  
plt.xlabel('sessionlength')  
plt.ylabel('Conversion (1 = Yes, 0 = No)')  
plt.show()
```



Group users based on VisitorType, OperatingSystems, and Region to identify potential differences in behavior and conversion rates.

```
[37]: grouped = shop_df.groupby(['VisitorType', 'OperatingSystems', 'Region'])
```

```
[38]: conversion_rate = grouped['Revenue'].mean() * 100
      conversion_rate.reset_index(name='ConversionRate')
```

```
[38]:
```

	VisitorType	OperatingSystems	Region	ConversionRate
0	New_Visitor	1	1	26.744186
1	New_Visitor	1	2	34.146341
2	New_Visitor	1	3	14.942529
3	New_Visitor	1	4	29.729730
4	New_Visitor	1	5	33.333333
..
110	Returning_Visitor	8	4	0.000000
111	Returning_Visitor	8	5	0.000000
112	Returning_Visitor	8	6	0.000000
113	Returning_Visitor	8	7	0.000000

```
114 Returning_Visitor          8          9          0.000000
```

```
[115 rows x 4 columns]
```

Segment users based on TrafficType and analyze their engagement patterns and purchase probability.

```
[39]: grouped = shop_df.groupby('TrafficType')
```

```
[40]: engagement_metrics = grouped[['PageValues', 'sessionlength']].mean()
      print(engagement_metrics)
```

	PageValues	sessionlength
TrafficType		
1	3.546226	35.078308
2	8.308613	41.774994
3	3.339503	28.428217
4	7.062934	31.458724
5	7.712489	21.842308
6	5.087703	32.164786
7	13.567345	32.600000
8	10.302436	29.498542
9	3.911694	17.414634
10	6.208230	35.751111
11	5.068642	27.344130
12	0.000000	3.000000
13	2.386929	35.571429
14	4.936097	85.076923
15	0.037454	18.378378
16	0.000000	18.333333
17	0.000000	4.000000
18	0.000000	16.300000
19	3.497520	41.352941
20	15.520252	23.036269


```
[41]: conversion_rate = grouped['Revenue'].mean() * 100
      print(conversion_rate)
```

TrafficType	
1	10.971524
2	21.656865
3	8.941878
4	15.478424
5	21.538462
6	11.963883
7	30.000000

```
8      27.696793
9       9.756098
10     20.000000
11     19.028340
12      0.000000
13      5.906593
14     15.384615
15      0.000000
16     33.333333
17      0.000000
18      0.000000
19      5.882353
20     25.906736
Name: Revenue, dtype: float64
```

```
[42]: engagement_metrics = grouped[['PageValues', 'sessionlength']].agg(['mean',  
    ↪ 'median'])
```

```
[43]: sns.barplot(x=engagement_metrics.index,  
    ↪ y=engagement_metrics['PageValues']['mean'])  
plt.title('Average PageValues by TrafficType')  
plt.xlabel('Traffic Type')  
plt.ylabel('Average PageValues')  
plt.show()
```



[43]:

[43]:

#Campaign Dataset

EDA on various features and columns.

[44]: `camp_df= pd.read_csv('/content/campaign - campaign.csv')`

[45]: `camp_df.head()`

[45]:

	ID	Year_Birth	Education	Marital_Status	Income	Kidhome	\
0	1826	1970	Graduation	Divorced	\$84,835.00	0	
1	1	1961	Graduation	Single	\$57,091.00	0	
2	10476	1958	Graduation	Married	\$67,267.00	0	
3	1386	1967	Graduation	Together	\$32,474.00	1	
4	5371	1989	Graduation	Single	\$21,474.00	1	

	Teenhome	Dt_Customer	Recency	MntWines	...	NumCatalogPurchases	\
0	0	6/16/14	0	189	...	4	
1	0	6/15/14	0	464	...	3	
2	1	5/13/14	0	134	...	2	
3	1	5/11/14	0	10	...	0	
4	0	4/8/14	0	6	...	1	

	NumStorePurchases	NumWebVisitsMonth	AcceptedCmp3	AcceptedCmp4	\
0	6	1	0	0	
1	7	5	0	0	
2	5	2	0	0	
3	2	7	0	0	
4	2	7	1	0	

	AcceptedCmp5	AcceptedCmp1	AcceptedCmp2	Complain	Country
0	0	0	0	0	SP
1	0	0	1	0	CA
2	0	0	0	0	US
3	0	0	0	0	AUS
4	0	0	0	0	SP

[5 rows x 27 columns]

```
[46]: camp_df.shape
```

```
[46]: (2239, 27)
```

```
[47]: camp_df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 2239 entries, 0 to 2238
Data columns (total 27 columns):
#   Column                Non-Null Count  Dtype
---  -
0   ID                    2239 non-null  int64
1   Year_Birth            2239 non-null  int64
2   Education             2239 non-null  object
3   Marital_Status        2239 non-null  object
4   Income                2239 non-null  object
5   Kidhome               2239 non-null  int64
6   Teenhome              2239 non-null  int64
7   Dt_Customer           2239 non-null  object
8   Recency               2239 non-null  int64
9   MntWines              2239 non-null  int64
10  MntFruits             2239 non-null  int64
11  MntMeatProducts       2239 non-null  int64
12  MntFishProducts       2239 non-null  int64
```

13	MntSweetProducts	2239	non-null	int64
14	MntGoldProds	2239	non-null	int64
15	NumDealsPurchases	2239	non-null	int64
16	NumWebPurchases	2239	non-null	int64
17	NumCatalogPurchases	2239	non-null	int64
18	NumStorePurchases	2239	non-null	int64
19	NumWebVisitsMonth	2239	non-null	int64
20	AcceptedCmp3	2239	non-null	int64
21	AcceptedCmp4	2239	non-null	int64
22	AcceptedCmp5	2239	non-null	int64
23	AcceptedCmp1	2239	non-null	int64
24	AcceptedCmp2	2239	non-null	int64
25	Complain	2239	non-null	int64
26	Country	2239	non-null	object

dtypes: int64(22), object(5)
memory usage: 472.4+ KB

```
[48]: camp_df['Dt_Customer'] = pd.to_datetime(camp_df['Dt_Customer'], format='%m/%d/%y')
```

```
[49]: camp_df['Income'] = camp_df['Income'].replace({'\$' : '' , ',' : ''} , regex=
↳ True ).astype(float)
```

```
[50]: camp_df.isnull().sum()
```

```
[50]: ID                                0
      Year_Birth                        0
      Education                        0
      Marital_Status                    0
      Income                            24
      Kidhome                           0
      Teenhome                          0
      Dt_Customer                       0
      Recency                           0
      MntWines                          0
      MntFruits                         0
      MntMeatProducts                   0
      MntFishProducts                   0
      MntSweetProducts                  0
      MntGoldProds                      0
      NumDealsPurchases                  0
      NumWebPurchases                    0
      NumCatalogPurchases               0
      NumStorePurchases                  0
      NumWebVisitsMonth                  0
      AcceptedCmp3                       0
      AcceptedCmp4                       0
```

```
AcceptedCmp5      0
AcceptedCmp1      0
AcceptedCmp2      0
Complain          0
Country           0
dtype: int64
```

```
[51]: camp_df[camp_df['Income'].isnull()]
```

```
[51]:
```

	ID	Year_Birth	Education	Marital_Status	Income	Kidhome	Teenhome	\
134	8996	1957	PhD	Married	NaN	2	1	
262	1994	1983	Graduation	Married	NaN	1	0	
394	3769	1972	PhD	Together	NaN	1	0	
449	5255	1986	Graduation	Single	NaN	1	0	
525	8268	1961	PhD	Married	NaN	0	1	
589	10629	1973	2n Cycle	Married	NaN	1	0	
898	10475	1970	Master	Together	NaN	0	1	
996	9235	1957	Graduation	Single	NaN	1	1	
1095	4345	1964	2n Cycle	Single	NaN	1	1	
1184	7187	1969	Master	Together	NaN	1	1	
1212	8720	1978	2n Cycle	Together	NaN	0	0	
1311	8557	1982	Graduation	Single	NaN	1	0	
1514	2863	1970	Graduation	Single	NaN	1	2	
1557	2437	1989	Graduation	Married	NaN	0	0	
1692	5250	1943	Master	Widow	NaN	0	0	
1803	7281	1959	PhD	Single	NaN	0	0	
1857	1612	1981	PhD	Single	NaN	1	0	
1862	5079	1971	Graduation	Married	NaN	1	1	
1879	10339	1954	Master	Together	NaN	0	1	
1966	5798	1973	Master	Together	NaN	0	0	
1982	2902	1958	Graduation	Together	NaN	1	1	
2138	3117	1955	Graduation	Single	NaN	0	1	
2164	7244	1951	Graduation	Single	NaN	2	1	
2169	1295	1963	Graduation	Married	NaN	0	1	

	Dt_Customer	Recency	MntWines	...	NumCatalogPurchases	\
134	2012-11-19	4	230	...	2	
262	2013-11-15	11	5	...	0	
394	2014-03-02	17	25	...	0	
449	2013-02-20	19	5	...	0	
525	2013-07-11	23	352	...	1	
589	2012-09-14	25	25	...	0	
898	2013-04-01	39	187	...	2	
996	2014-05-27	45	7	...	0	
1095	2014-01-12	49	5	...	0	
1184	2013-05-18	52	375	...	10	
1212	2012-08-12	53	32	...	0	

1311	2013-06-17	57	11	...	0
1514	2013-08-23	67	738	...	3
1557	2013-06-03	69	861	...	5
1692	2013-10-30	75	532	...	5
1803	2013-11-05	80	81	...	3
1857	2013-05-31	82	23	...	0
1862	2013-03-03	82	71	...	1
1879	2013-06-23	83	161	...	1
1966	2013-11-23	87	445	...	4
1982	2012-09-03	87	19	...	0
2138	2013-10-18	95	264	...	1
2164	2014-01-01	96	48	...	1
2169	2013-08-11	96	231	...	5

	NumStorePurchases	NumWebVisitsMonth	AcceptedCmp3	AcceptedCmp4	\
134	8	9	0	0	
262	2	7	0	0	
394	3	7	0	0	
449	0	1	0	0	
525	7	6	0	0	
589	3	8	0	0	
898	6	5	0	0	
996	2	7	0	0	
1095	2	7	0	0	
1184	4	3	0	0	
1212	1	0	0	1	
1311	3	6	0	0	
1514	10	7	0	1	
1557	12	3	0	1	
1692	11	1	0	0	
1803	4	2	0	0	
1857	3	6	0	0	
1862	3	8	0	0	
1879	4	6	0	0	
1966	8	1	0	0	
1982	3	5	0	0	
2138	5	7	0	0	
2164	4	6	0	0	
2169	7	4	0	0	

	AcceptedCmp5	AcceptedCmp1	AcceptedCmp2	Complain	Country
134	0	0	0	0	GER
262	0	0	0	0	US
394	0	0	0	0	AUS
449	0	0	0	0	AUS
525	0	0	0	0	CA
589	0	0	0	0	GER

898	0	0	0	0	US
996	0	0	0	0	GER
1095	0	0	0	0	AUS
1184	0	0	0	0	AUS
1212	0	0	0	0	IND
1311	0	0	0	0	AUS
1514	0	1	0	0	SP
1557	0	1	0	0	SP
1692	1	0	0	0	AUS
1803	0	0	0	0	AUS
1857	0	0	0	0	AUS
1862	0	0	0	0	AUS
1879	0	0	0	0	AUS
1966	0	0	0	0	GER
1982	0	0	0	0	AUS
2138	0	0	0	0	AUS
2164	0	0	0	0	AUS
2169	0	0	0	0	CA

[24 rows x 27 columns]

```
[52]: camp_df['Income'] = camp_df['Income'].fillna(0)
```

```
[53]: camp_df['Income'].isnull().sum()
```

```
[53]: 0
```

Feature Engineering

```
[54]: camp_df['Customer_Lifetime'] = (pd.to_datetime('today') -
    ↪ camp_df['Dt_Customer']).dt.days
```

- Binning

```
[55]: camp_df['Recency_Binned'] = pd.cut(camp_df['Recency'], bins=[0, 30, 60, 90,
    ↪ 120, 180], labels=['Very Recent', 'Recent', 'Moderate', 'Old', 'Very Old'])
```

```
[56]: camp_df['Recency_Binned'].value_counts()
```

```
[56]: Recency_Binned
Very Recent    695
Moderate       664
Recent         654
Old            198
Very Old         0
Name: count, dtype: int64
```


Remove extreme values if required.



```
[57]: numerical_cols = camp_df.select_dtypes(include=np.number).columns
```



```
[58]: numerical_cols
```



```
[58]: Index(['ID', 'Year_Birth', 'Income', 'Kidhome', 'Teenhome', 'Recency',  
         'MntWines', 'MntFruits', 'MntMeatProducts', 'MntFishProducts',  
         'MntSweetProducts', 'MntGoldProds', 'NumDealsPurchases',  
         'NumWebPurchases', 'NumCatalogPurchases', 'NumStorePurchases',  
         'NumWebVisitsMonth', 'AcceptedCmp3', 'AcceptedCmp4', 'AcceptedCmp5',  
         'AcceptedCmp1', 'AcceptedCmp2', 'Complain', 'Customer_Lifetime'],  
        dtype='object')
```



```
[59]: for col in numerical_cols:  
    plt.figure(figsize=(10,5))  
    sns.boxplot(camp_df[col])  
    plt.show()
```










































```
[60]: key_columns = [
        'Income', 'MntWines', 'MntFruits', 'MntMeatProducts', 'MntFishProducts',
        'MntSweetProducts', 'MntGoldProds'
    ]
```

```
[61]: Q1 = camp_df[key_columns].quantile(0.25)
      Q3 = camp_df[key_columns].quantile(0.75)
      IQR = Q3 - Q1

      lower_bound = Q1 - 1.5 * IQR
      upper_bound = Q3 + 1.5 * IQR

      outliers = ((camp_df[key_columns] < lower_bound) | (camp_df[key_columns] >
      ↪upper_bound)).any(axis=1)

      camp_df[outliers]
```

```
[61]:
```

	ID	Year_Birth	Education	Marital_Status	Income	Kidhome	\
0	1826	1970	Graduation	Divorced	84835.0	0	
5	7348	1958	PhD	Single	71691.0	0	
10	2079	1947	2n Cycle	Married	81044.0	0	
12	10530	1959	PhD	Widow	67786.0	0	
14	10311	1969	Graduation	Married	4428.0	0	
...	
2218	5687	1980	Graduation	Divorced	81702.0	0	
2222	2831	1976	Graduation	Together	78416.0	0	

2225	1743	1974	Graduation	Single	69719.0	0
2237	528	1978	Graduation	Married	65819.0	0
2238	4070	1969	PhD	Married	94871.0	0

	Teenhome	Dt_Customer	Recency	MntWines	...	NumWebVisitsMonth	\
0	0	2014-06-16	0	189	...	1	
5	0	2014-03-17	0	336	...	2	
10	0	2013-12-27	0	450	...	1	
12	0	2013-12-07	0	431	...	1	
14	1	2013-10-05	0	16	...	1	
...	
2218	0	2012-09-23	98	563	...	3	
2222	1	2014-06-27	99	453	...	3	
2225	0	2014-05-26	99	273	...	1	
2237	0	2012-11-29	99	267	...	3	
2238	2	2012-09-01	99	169	...	7	

	AcceptedCmp3	AcceptedCmp4	AcceptedCmp5	AcceptedCmp1	AcceptedCmp2	\
0	0	0	0	0	0	
5	0	0	0	0	0	
10	0	0	0	0	0	
12	0	0	0	0	0	
14	0	0	0	0	0	
...	
2218	0	0	0	0	0	
2222	0	0	0	0	0	
2225	0	0	0	0	0	
2237	0	0	0	0	0	
2238	0	1	1	0	0	

	Complain	Country	Customer_Lifetime	Recency_Binned
0	0	SP	3727	NaN
5	0	SP	3818	NaN
10	0	US	3898	NaN
12	0	IND	3918	NaN
14	0	SP	3981	NaN
...
2218	0	CA	4358	Old
2222	0	SP	3716	Old
2225	0	SP	3748	Old
2237	0	IND	4291	Old
2238	0	CA	4380	Old

[630 rows x 29 columns]

```
[62]: for col in key_columns:
        median = camp_df[col].median()
```

```
camp_df.loc[outliers, col] = median
```


Hypothesis testing:


- Is income of customers dependent on their education


One-way Anova


Assumptions


```
[63]: for level in camp_df['Education'].unique():  
      income_data = camp_df[camp_df['Education'] == level]['Income']  
      stats.probplot(income_data, dist="norm", plot=plt)  
      plt.title(f'Q-Q Plot for {level}')  
      plt.show()
```












```
[64]: sns.boxplot(x='Education', y='Income', data=camp_df)
plt.title('Box Plot of Income by Education Level')
plt.xlabel('Education Level')
plt.ylabel('Income')
plt.xticks(rotation=45)
plt.show()
```

```
[65]: grouped = camp_df.groupby('Education')['Income']
```

```
[66]: Graduation = grouped.get_group('Graduation')
      PhD = grouped.get_group('PhD')
      Master = grouped.get_group('Master')
      Basic = grouped.get_group('Basic')
      twon_Year = grouped.get_group('2n Cycle')
```

```
[67]: anova_result = stats.f_oneway(Graduation, PhD, Master, Basic, twon_Year)
      anova_result
```

```
[67]: F_onewayResult(statistic=52.71479834410078, pvalue=1.7164204702604772e-42)
```

```
[68]: alpha = 0.05
      if anova_result.pvalue < alpha:
          print("Reject the null hypothesis. There is a significant difference in_
          income between education levels.")
```

```

else:
    print("Fail to reject the null hypothesis. There is no significant_
    ↪difference in income between education levels.")

```

Reject the null hypothesis. There is a significant difference in income between education levels.

Do higher income people spend more (take in account spending in all categories together)

```

[69]: camp_df['Total_Spending'] = camp_df['MntWines'] + camp_df['MntFruits'] +_
    ↪camp_df['MntMeatProducts'] + camp_df['MntFishProducts'] +_
    ↪camp_df['MntSweetProducts']

```

```

[70]: corr , p_value = stats.pearsonr(camp_df['Income'], camp_df['Total_Spending'])

```

```

[71]: corr, p_value

```

```

[71]: (0.7091856965640443, 0.0)

```

```

[72]: alpha = 0.05
    if p_value < alpha:
        print("Reject the null hypothesis. There is a significant correlation_
        ↪between income and total spending.")
    else:
        print("Fail to reject the null hypothesis. There is no significant_
        ↪correlation between income and total spending.")

```

Reject the null hypothesis. There is a significant correlation between income and total spending.

Do couples spend more or less money on wine than people living alone (set 'Married', 'Together': 'In couple' and 'Divorced', 'Single', 'Absurd', 'Widow', 'YOLO': 'Alone')

```

[73]: camp_df['Living_Status'] = camp_df['Marital_Status'].apply(lambda x: 'Couple'_
    ↪if x in ['Married', 'Together'] else 'Alone')

```

```

[74]: camp_df['Living_Status'].value_counts()

```

```

[74]: Living_Status
Couple    1443
Alone      796
Name: count, dtype: int64

```

```

[75]: wine_spending_couple = camp_df[camp_df['Living_Status'] == 'Couple']['MntWines']
    wine_spending_alone = camp_df[camp_df['Living_Status'] == 'Alone']['MntWines']

```

```
[76]: h_stat, p_value = stats.kruskal(wine_spending_couple, wine_spending_alone)
```

```
[77]: h_stat, p_value
```

```
[77]: (0.5516579031462469, 0.45764106602403265)
```

```
[78]: alpha = 0.05
      if p_value < alpha:
          print("Reject the null hypothesis. There is a significant difference in_
              ↪ wine spending between couples and people living alone.")
      else:
          print("Fail to reject the null hypothesis. There is no significant_
              ↪ difference in wine spending between couples and people living alone.")
```

Fail to reject the null hypothesis. There is no significant difference in wine spending between couples and people living alone.

Are people with lower income are more attracted towards campaign or simply put accept more campaigns. (create two income brackets one below median , other above median income and create a column which tells if they have ever accepted any campaign)

```
[79]: medium_income = camp_df['Income'].median()
```

```
[80]: camp_df['Income_Category'] = camp_df['Income'].apply(lambda x: 'Low Income' if_
              ↪ x < medium_income else 'High Income')
```

```
[83]: camp_df['Total_Acceptency'] = camp_df['AcceptedCmp1'] + camp_df['AcceptedCmp2']_
              ↪ + camp_df['AcceptedCmp3'] + camp_df['AcceptedCmp4'] + camp_df['AcceptedCmp5']
```

```
[84]: camp_df['Accepted_any_campaign'] = camp_df.groupby('ID')['Total_Acceptency'].
              ↪ transform(lambda x: 1 if x.sum() > 0 else 0 )
```

```
[85]: acceptance_rate = camp_df.groupby('Income_Category')['Accepted_any_campaign'].
              ↪ mean() * 100
```

```
[86]: acceptance_rate
```

```
[86]: Income_Category
      High Income      30.901288
      Low Income       9.590317
      Name: Accepted_any_campaign, dtype: float64
```

Approximately 31% of high-income individuals have accepted at least one campaign.

This suggests that individuals with higher income are more likely to accept campaigns. This could be because they have more disposable income, feel more targeted by the campaigns, or simply

because the campaigns are better tailored to their preferences.

[]: